

# CONSIDERATIONS FOR INTEGRATING VIRTUAL THREADS IN A JAVA FRAMEWORK: A QUARKUS EXAMPLE IN A RESOURCE-CONSTRAINED ENVIRONMENT

A. NAVARRO<sup>1,2</sup> J. PONGE<sup>1,2</sup> F. LE MOUËL<sup>2</sup> C. ESCOFFIER<sup>1</sup>

<sup>1</sup>RED HAT

<sup>2</sup>CITI LABORATORY-INSA LYON



*17TH ACM INTERNATIONAL CONFERENCE ON DISTRIBUTED AND  
EVENT-BASED SYSTEMS 2023*

**1** Context

2 Integration

3 The experiment

4 Results



MyApplication


# THE END OF THE MONOLITH

- Difficult to maintain
- Dependency hell
- Reboot for every change
- Sub-optimal deployment
- Limit scalability
- Technology lock-in

from

Microservices: yesterday, today, and tomorrow - Dragoni & al

# MICROSERVICES



# MICROSERVICES

- *Difficult to maintain*  
Split in smaller entities
- *Dependency hell*  
Less dependencies per service
- *Reboot for every change*  
Reboot only impacted service
- *Sub-optimal deployment*  
Per-service deployment
- *Limits scalability*  
Scale each service
- *Technology lock-in*  
Per-service technology

## Remark

Transitioning to cloud becomes cost efficient if done well.

# MAKE MICROSERVICES CHEAP

resource-efficiency

# MAKE MICROSERVICES CHEAP

resource-efficiency  
resource-efficiency



# MAKE MICROSERVICES CHEAP

resource-efficiency

resource-efficiency

resource-efficiency

# MAKE MICROSERVICES CHEAP

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

# MAKE MICROSERVICES CHEAP

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

resource-efficiency

# MAKE MICROSERVICES CHEAP

resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency

resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency

resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency  
resource-efficiency

# SCALING AND REACTIVE PROGRAMMING

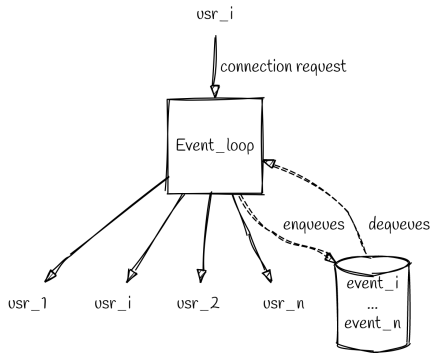

## Problem

Communications over the network are unreliable and slow.

## Goal

Achieving efficient and timely communication between services.

# BLOCKING AND NON-BLOCKING PROGRAMMING



## CODE COMPARISON - BLOCKING

---

```
var names = getAll();
var quotes = getQuotes(names.size());
for(int i=0; i < names.size();i++){
    names.get(i).surname+= "- "+quotes.get(i);
}
return names;
```

---

## CODE COMPARISON - ASYNCHRONOUS CALLBACKS

---

```
getAll( names => {  
    getQuotes(names.size(), quotes => {  
        for(int i=0; i < names.size();i ++){  
            names.get(i).surname+= "- "+quotes.get(i);  
        }  
        //continuation  
    })  
});
```

---



## CODE COMPARISON - REACTIVE STREAMS

---

```
var names = getAll().memoize().indefinitely();
var quotes = names.onItem().transformToUni(list ->
    getQuotes(list.size()));
return Uni.combine().all()
    .unis(names,quotes).asTuple()
    .onItem().transform(tuple -> {
        var nList=tuple.getItem1();
        //can await it since it is already resolved
        var qList = tuple.getItem2();
        for(int i=0; i < namesList.size();i++){
            nList.get(i).surname += " - "+qList.get(i);
        }
        return namesList;
    });
```

---

*We should do (as wise programmers aware of our limitations) our utmost best to shorten the conceptual gap between the static program and the dynamic process, to **make the correspondence between the program (spread out in text space) and the process (spread out in time) as trivial as possible.***

*Go To Statement considered harmful.* - Edgar. J. Dijkstra, 1968


## CODE COMPARISON - VIRTUAL THREADS

---

```
var names = getAll();
var quotes = getQuotes(names.size());
for(int i=0; i < names.size();i++){
    names.get(i).surname+= "- "+quotes.get(i);
}
return names;
```

---

# VIRTUAL THREADS



# TABLE OF CONTENTS

1 Context

**2 Integration**

3 The experiment

4 Results

# A QUICK OVERVIEW OF QUARKUS

---

```
@GET
@Path("/named_quotes")
Uni<List<Name>> myEndpoint() {
    var names = getAll().memoize().indefinitely();
    var quotes = names.onItem().transformToUni(list ->
        getQuotes(list.size()));
    return Uni.combine().all()
        .unis(names,quotes).asTuple()
        .onItem().transform(tuple -> {
            var nList=tuple.getItem1();
            var qList = tuple.getItem2();
            for(int i=0; i < namesList.size();i ++){
                nList.get(i).surname += " - "
                    +qList.get(i);
            }
            return namesList;
        });
}
```

# A QUICK OVERVIEW OF QUARKUS-VIRTUAL-THREAD

---

```
@GET
@Path("/named_quotes")
@RunOnVirtualThread
List<Name> myBetterEndpoint() {
    var names = getAll();
    var quotes = getQuotes(names.size());
    for(int i=0; i < names.size(); i++){
        names.get(i).surname+= "- "+quotes.get(i);
    }
    return names;
}
```

---

## THREE INTEGRATION STRATEGIES



Forking the worker model.

Using event-loops as carriers.


“Virtualizing” Netty event-loops.





# FORKING THE WORKER MODEL



# REUSING EVENT-LOOPS



# "VIRTUALIZING" NETTY EVENT-LOOPS





# INTEGRATION STRATEGIES SUMMARY

| Strategy  | Pros  | Cons   |
|---|---|--|
| Forking worker model                              | Simple, fits virtual threads model  | Context switches                                   |
| Using event-loop as carrier                       | No context-switch, Fewer threads overall  | Potential deadlocks                                |
| Modifying Netty event-loops to be virtual threads | Integration done at the Netty level, Netty-based frameworks would benefit from it | Can't modify Netty upstream, unpredictable effects |

**Table:** Comparison of the different Quarkus-virtual-threads options

# FINAL CHOICE: FORKING THE WORKER MODEL



# TABLE OF CONTENTS

1 Context

2 Integration

**3 The experiment**

4 Results

# THE EXPERIMENT

## Goal

Measure how performance of the application is affected by replacing reactive endpoints with virtual-threads offloading

## Hypothesis

**Quarkus-virtual-threads** should perform better than **Quarkus-blocking** but not as well as **Quarkus-reactive**

## Limited resources


- 512MB memory
- 0.5 vCPU
- 256MB heap

## Fault-inducing settings

- 200ms of delay between the DB and the server
- Hyperfoil to avoid Coordinated Omission



# ELEMENTS OF THE EXPERIMENT



# TABLE OF CONTENTS





1 Context

2 Integration

3 The experiment

**4 Results**

# RESULTS WITH 200MS DELAY



# GARBAGE COLLECTION DETAIL

|                       | Quarkus-virt-threads-0 | Quarkus-virt-threads-200 |
|-----------------------|------------------------|--------------------------|
| Max-latency           | 1.44 s                 | 31.54 s                  |
| GC count              | 275                    | 709                      |
| Avg pause             | 18.662 ms              | 92.270 ms                |
| Longest pause         | 89.723 ms              | 520.437 ms               |
| Young Collection time | 5.132 s                | 41.685 s                 |
| Old Collection time   | N/A                    | 23.734 s                 |
| Sum of pauses         | 5.132 s                | 65 s                     |

|                       | Quarkus-reactive-0 | Quarkus-reactive-200 |
|-----------------------|--------------------|----------------------|
| Max-latency           | 1.08 s             | 704.64 ms            |
| GC count              | 192                | 183                  |
| Avg pause             | 15.169 ms          | 14.968 ms            |
| Longest pause         | 32.898 ms          | 49.312 ms            |
| Young Collection time | 3.004              | 2.739s               |
| Old Collection time   | N/A                | N/A                  |
| Sum of pauses         | 3.004 s            | 2.739 s              |

# EXPLAINING LONGER GARBAGE COLLECTION

- Too many virtual threads exist in memory at the same time.
- Data structures generated by Quarkus/Netty pollute memory.

# EXPLAINING LONGER GARBAGE COLLECTION

- Too many virtual threads exist in memory at the same time.
- Data structures generated by Quarkus/Netty pollute memory.

## GENERALIZING THE RESULTS

*The integration of virtual threads in frameworks built on top of Netty or that rely heavily on ThreadLocals, presuming a small thread count, is likely to experience significant garbage collection (GC) pressure.*

# CONCLUSIONS

- *Quarkus-blocking* < *Quarkus-virtual-threads* < *Quarkus-reactive*.
- High Memory consumption.
- Virtual threads are not responsible for the high memory consumption.
- Data structures generated by *Quarkus/Netty* are the main memory consumers.


*arnavarr@redhat.com*



# DEADLOCK SITUATION

```
@GET
@RunOnVirtualThread
@NonBlocking
@Path("/print-both")
public List<Fortune> getAll() {
    System.out.println("outer - " + Thread.currentThread());
    var list =
        db.getPool().preparedQuery(SELECT_ALL)
            .execute()
            .map(item -> {
                System.out.println("inner - " + Thread.currentThread());
                return createListOfFortunes(item);
            });
    return list.await().indefinitely();
}
```

# DEADLOCK EXPLANATION



## Conclusion

The event-loop can't reuse locks *as a carrier*

# RESULTS WITHOUT DELAY

